



ML@B Bootcamp - Basic Models in Tensorflow

Machine Learning at Berkeley

Agenda

Basic TensorFlow Structure

Lab Overview

Basic TensorFlow Structure

Loading in Datasets isn't terribly challenging and we'll demonstrate with a csv.

Many available general interfaces for loading

Tensorflow has a nice interface for MNIST that provides many convenient methods

Placeholders act as the "input" markers for our TF graph

We need to define placeholders for our features X and the corresponding labels, y .

The first dimension of X and y placeholders are always set to `None` - this signals to the graph constructor that we want to accept variable batch sizes to our model.

Placeholders not just limited to training data - sometimes we can use them for other values like dropout probability.

Step 3: Define Variables for Each Layer



Variables hold the parameters of our model.

Fully-connected and Convolutional layers: both use weight and bias variables.

Weights - typically initialized from a normal random distribution to avoid local minimum and break symmetry in backprop that would occur with zero initialization

Biases - typically initialized as slightly positive (e.g. 0.1)

Step 4: Construct each layer



We limit our focus to fully-connected layers for now

We need to construct the following function for each layer:

$$\sigma(X \cdot W + b)$$

- σ will be the activation function of the layer (relu, sigmoid)
- X is the input from the previous layer (or the input data for the first layer). The size will [batch size, num. input features.]
- W are the weights, and will be of size [num. input features, num. output features]
- b is the bias, and will be of size [1, num. output features]

Loss fn simplicity:

- Define our objective
- Function that takes in the Y placeholder and compares to the net prediction, \hat{Y}
- Eg:
 - Simple : Mean Squared Error $\mathcal{L} = \|Y - \hat{Y}\|_2^2$
 - More advanced: Cross Entropy
$$H_{y'}(y) := - \sum_i (y'_i \log(y_i) + (1 - y'_i) \log(1 - y_i))$$
 - Super advanced: Adding other quantities to loss functions (eg Style transfer with style loss and content loss)
 - We are provided fns like [cross entropy](#)

- Takes in learning rate as parameter
- If loss is failing to converge with each epoch of training or appearing as nan, try decreasing the learning rate
- Set the optimizer to minimize the loss function

- Can use `InteractiveSession()` for iPython notebooks and `Session()` for Python scripts
- Make sure to initialize all global variables using:
`sess.run(tf.global_variables_initializer())`

- For each epoch, go through all batches of training data and run optimizer
- Aggregate loss over all batches in each epoch - print this value at the end of the epoch
- For models that take longer to train, may be a good idea to create a function that runs the optimization process for a given number of epochs
 - Allows you to see if loss is decreasing after certain number of epochs

- Depends on type of problem you are trying to solve
- Can include printing accuracy, plotting a regression line, displaying generated data, etc.

Lab Overview

- Linear Regression on Dummy Data
- Logistic Regression on MNIST Data
- Basic Feed Forward Network on MNIST Data

Extra Exercise - Denoising Autoencoder

