



ML@B Bootcamp - Convolutional Neural Networks in Tensorflow

Machine Learning at Berkeley

Agenda



Why Use CNNs?

CNN Layers


Tensorflow Syntax

Transfer Learning

Why Use CNNs?



- Spatial information is lost if image is converted to vector




- CNNs arrange each layer of neurons into 3 dimensions




- A 200x200 color image has $200 \times 200 \times 3 = 120,000$ pixels
- Fully connected network will have 120,000 weights in first layer
- Having such huge number of parameters is wasteful for memory and training, and leads to overfitting


CNN Layers




- Suppose input image has dimension $32 \times 32 \times 3$, and each of 5 filters has dimension $5 \times 5 \times 3$
- Convolution operation is computing dot product of each filter along each position of the image



- Each filter only looks at small region of image but applies over all channels
- Multiple filters look at same region of image



- **Depth of Output Volume** - number of filters used on each input layer



- **Stride** - how many pixels in each dimension to move the filter

0	0	0	0	0	0			
0								
0								
0								
0								

- **Zero Padding** - helps control height and width of output volume

- If a square input layer has dimensions $W_1 \times W_1 \times D_1$ and N square filters of dimensions $F \times F \times D_1$ are applied with zero padding P and stride S , output volume $W_2 \times W_2 \times D_2$ is given by:

$$W_2 = \frac{W_1 - F + 2P}{S} + 1$$

$$D_2 = N$$


- Eg: If input volume has size $32 \times 32 \times 3$ and 10 5×5 filters are applied with stride 1 and zero pad 2, what is output volume?

- If a square input layer has dimensions $W_1 \times W_1 \times D_1$ and N square filters of dimensions $F \times F \times D_1$ are applied with zero padding P and stride S , output volume $W_2 \times W_2 \times D_2$ is given by:

$$W_2 = \frac{W_1 - F + 2P}{S} + 1$$


$$D_2 = N$$


- Eg: If input volume has size $32 \times 32 \times 3$ and 10 5×5 filters are applied with stride 1 and zero pad 2, what is output volume?
- $W_2 = \frac{32 - 5 + 2 \times 2}{1} + 1 = 32$ so output volume is $32 \times 32 \times 10$



- Activation function which adds nonlinearity

- Reduces size of representation
- Applied on each channel independently instead of acting over depth of input





- Takes in input of size $W_1 \times W_1 \times D_1$ and applies filter of size $F \times F$ with stride S to produce output of size $W_2 \times W_2 \times D_2$ where $D_2 = D_1$ and $W_2 = \frac{W_1 - F}{S} + 1$

Tensorflow Syntax

tf.nn.conv2d

```
conv2d(  
    input,  
    filter,  
    strides,  
    padding,  
    use_cudnn_on_gpu=None,  
    data_format=None,  
    name=None  
)
```

- input: 4d tensor of shape [number of examples in batch, height of input, width of input, number of channels of input]
- filter: 4d tensor of shape [filter height, filter width, number of input channels, number of output channels]
- Note: number of output channels = number of filters used

tf.nn.conv2d

```
conv2d(  
    input,  
    filter,  
    strides,  
    padding,  
    use_cudnn_on_gpu=None,  
    data_format=None,  
    name=None  
)
```

- strides: 1D tensor of length 4 of format [1, stride vertical, stride horizontal, 1]
- padding: "SAME" or "VALID"

```
tf.layers.conv2d(inputs, filters, kernel_size, strides=(1, 1), padding='valid',  
data_format='channels_last', dilation_rate=(1, 1), activation=None,  
use_bias=True, kernel_initializer=None,  
bias_initializer=tf.zeros_initializer(), kernel_regularizer=None,  
bias_regularizer=None, activity_regularizer=None, trainable=True, name=None,  
reuse=None)
```

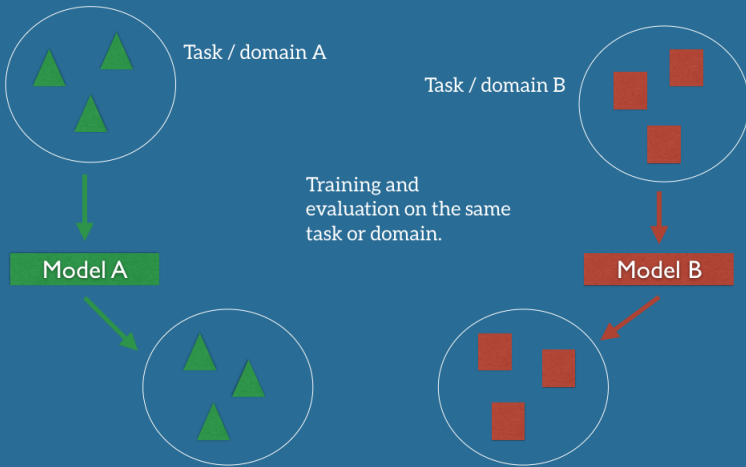
- input: 4d tensor of shape [number of examples in batch, height of input, width of input, number of channels of input]
- filters: integer of number of filters in convolution or depth of output space
- kernel_size: integer or tuple specifying width and height of filter

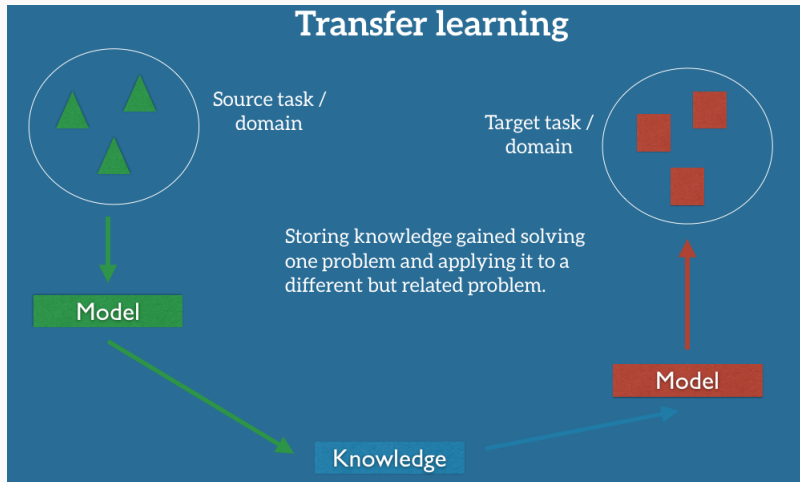
- weight tensor must be initialized in `tf.nn.conv2d` whereas `tf.layers.conv2d` only needs size of filter and output
- `tf.layers.conv2d` may be easier when creating and training model from scratch
- `tf.nn.conv2d` may be easier when loading pretrained model

Transfer Learning

- Transfer learning takes the knowledge of pretrained networks and adapts them to new applications
- For example, if you want to classify domain-specific objects (say different species of birds) in images you might want to start with VGG16 (trained on imagenet)
- Imagenet isn't focused on species of birds (and thus VGG16) isn't focused on birds, but it's (much) better than nothing


Traditional ML





- Don't have to reinvent the wheel
- Is a lot faster than training a whole new model from scratch
- Don't have to gather as much data

How to Actually Transfer Learn



- Freeze the convolutional layers with imagenet weights
- Train the "top" three FC layers

- Transfer learning can also be viewed as training a "feature extractor"
- For example, VGG16 takes a 224×224 image and outputs a 512×1 feature vector

- Goal: Classify cats vs dogs
- Take a pretrained convnet
- Remove the last fully-connected layer
- Extract the features
- Train a classifier

- Data Augmentation: rotation, shearing, random cropping, flipping, stretching

```
train_datagen = ImageDataGenerator(  
    preprocessing_function=preprocess_input,  
    rotation_range=30,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True  
)  
  
train_generator = train_datagen.flow_from_directory(  
    args.train_dir,  
    target_size=(IM_WIDTH, IM_HEIGHT),  
    batch_size=batch_size,  
)
```

- Initialize Inception V3 network from `keras.applications` module
- Add last layer to convnet
 - Input is output of base model
 - Add pooling layer with `GlobalAveragePooling2D()`
 - Add fully connected `Dense()` layer of size 1024 with a softmax function
 - Return new model with input base model and output last layer
- Freeze all layers and compile the new model

- If new dataset is small and similar to original \Rightarrow Retrain final linear classifier
- If new dataset is large and similar to original \Rightarrow Fine tune through whole network
- If new dataset is small and different from original \Rightarrow Retrain later layers in network
- If new dataset is large and different from original \Rightarrow Use weights of previous network to initialize new network

- Replace and retrain classifier on top of convnet
- Note: Pretrained network's weights are tuned
- When to fine tune

	Similar dataset	Different dataset
Small dataset	Transfer learning: highest level features + classifier	Transfer learning: lower level features + classifier
Large dataset	Fine-tune*	Fine-tune*

- Freeze all layers and compile model
 - Set each layer.trainable = False
 - Compile model with rmsprop optimizer, categorical cross entropy loss, and accuracy metrics
 - Lower learning rate from learning rate used during training from scratch

- Initialize Inception V3 network from keras.applications module

```
base_model = InceptionV3(weights='imagenet', include_top=False)
```

- Add last layer to convnet

```
def add_new_last_layer(base_model, nb_classes):
    """Add last layer to the convnet
    Args:
        base_model: keras model excluding top
        nb_classes: # of classes
    Returns:
        new keras model with last layer
    """
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(FC_SIZE, activation='relu')(x)
    predictions = Dense(nb_classes, activation='softmax')(x)
    model = Model(input=base_model.input, output=predictions)
    return model
```

- Freeze all layers and compile the new model

```
def setup_to_transfer_learn(model, base_model):
    """Freeze all layers and compile the model"""
    for layer in base_model.layers:
        layer.trainable = False
    model.compile(optimizer='rmsprop',
```

- Freeze all layers and compile the new model

```
def setup_to_transfer_learn(model, base_model):
    """Freeze all layers and compile the model"""
    for layer in base_model.layers:
        layer.trainable = False
    model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```

- Finetuning: Freeze all layers and compile model

```
def setup_to_finetune(model):
    """Freeze the bottom NB_IV3_LAYERS and retrain the remaining top
    layers.
    note: NB_IV3_LAYERS corresponds to the top 2 inception blocks in
    the inceptionv3 architecture
    Args:
        model: keras model
    """
    for layer in model.layers[:NB_IV3_LAYERS_TO_FREEZE]:
        layer.trainable = False
    for layer in model.layers[NB_IV3_LAYERS_TO_FREEZE:]:
        layer.trainable = True
    model.compile(optimizer=SGD(lr=0.0001, momentum=0.9),
                  loss='categorical_crossentropy')
```

```
python predict.py --image dog.001.jpg --model dc.model  
python predict.py --image_url https://goo.gl/Xws7Tp --model dc.model
```

